UNITED INTERNATIONAL COLLEGE

REGRESSION ANALYSIS USING R

Beijing PM2.5 Data Analysis

Author:
Junjie LIU
Daichen YAO
Yun YANG

Supervisor:
Dr.Ye HUA JUN
Dr.He PING

Group S
Division of Science and Technology
Statistics

December 13, 2018

Declaration of Authorship

I, Junjie LIU
Daichen YAO
Yun YANG, declare that this thesis titled, "Beijing PM2.5 Data Analysis" and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed: Junjie LIU Daichen YAO Yun YANG

Date: December 13, 2018

Contents

List of Figures

Chapter 1

Introduction

1.1 Data Set Selection

In the very beginning, we have been shown the data set website UCI Data Sets, we selected the **"Beijing PM2.5 Data Data Set**" which was donated by Prof.Song Xi Chen, from Guanghua School of Management, Center for Statistical Science, Peking University(Liang et al., 2015), for Beijing is known around the world by severe air pollution. The most widely used method to measure air pollution is PM2.5, the particulate matter smaller than or equal to 2.5 microns in diameter in the air, which is a kind of air pollutant that people and our daily life pay great attention to. Although this data set is time-dependent, as long as the data pre-processing is reasonable, we can still treat the data at each time point as independent and unrelated data for regression analysis.

1.2 Data Sets Introduction

In the original data sets, we can found that it contains almost every hour's PM2.5, Dew Point, Temperature, Pressure, Combined wind direction, Cumulative wind speed, Cumulative hours of snow and Cumulative hours of rain data, and in this project, our purposes are to determine the major factors responsible for this pollution, discuss whether these factors have been targeted by recent initiatives, and predict future PM2.5 levels.

1.3 Multiple Regression Estimation

In statistical modeling, regression model is one of the most important models. It is estimating the relationship between dependent variable and independent variables by the observation data. Also, the parameters are unbiased.(Montgomery, Peck, and Vining, 2012).

Regression model involve the following parameters and variables:

- **unknown parameters**, denoted as β, which may represented as a vector;
- **independent variables**, denoted as $\mathbf{X}$, which is represented as a matrix;
- **dependent variable**, denoted as $\mathbf{Y}$, which is represented as a vector.

The **Ordinary Multiple Regression Model** is

$$\mathbf{Y} = \mathbf{X}\beta + \epsilon$$

We think this kind of estimation may suitable for our data.

Chapter 2

Data Pre-processing

2.1 Missing Values

At the very beginning, the figure below shows the head of the original data set.

	No	year	month	day	hour	pm2.5	DEWP	TEMP	PRES	cbwd	Iws	Is	Ir
	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>	<dbl>	<fct>	<dbl>	<int>	<int>
1	1	2010	1	1	0	NA	-21	-11	1021	NW	1.79	0	0
2	2	2010	1	1	1	NA	-21	-12	1020	NW	4.92	0	0
3	3	2010	1	1	2	NA	-21	-11	1019	NW	6.71	0	0
4	4	2010	1	1	3	NA	-21	-14	1019	NW	9.84	0	0
5	5	2010	1	1	4	NA	-20	-12	1018	NW	13.0	0	0
6	6	2010	1	1	5	NA	-19	-10	1017	NW	16.1	0	0
7	7	2010	1	1	6	NA	-19	-9	1017	NW	19.2	0	0
8	8	2010	1	1	7	NA	-19	-9	1017	NW	21.0	0	0
9	9	2010	1	1	8	NA	-19	-9	1017	NW	24.2	0	0
10	10	2010	1	1	9	NA	-20	-8	1017	NW	27.3	0	0

... with 43,814 more rows

FIGURE 2.1: Origin Data Set

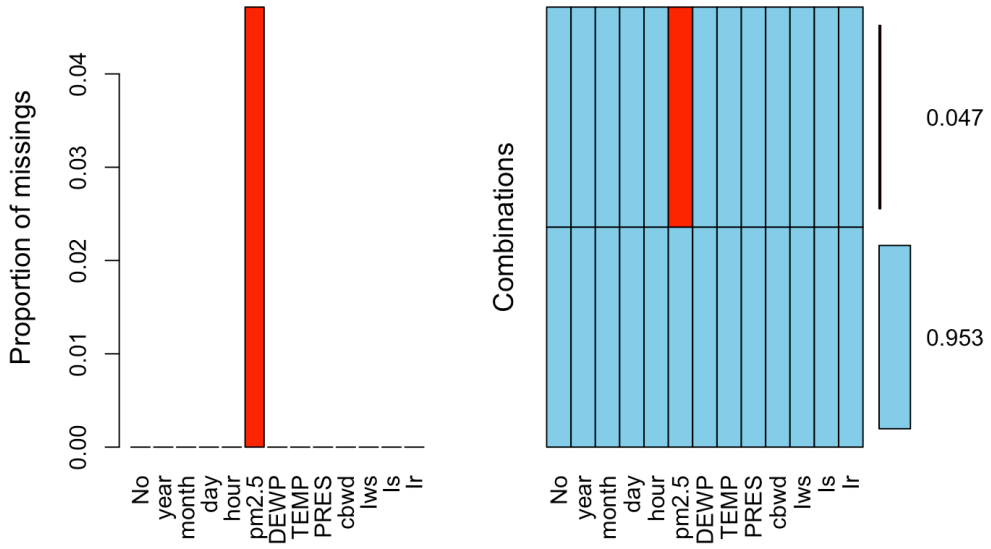

FIGURE 2.2: Visualization of Missing Value

The data set we chosen is from the UCI Machine Learning Repository datasets, it has collected Beijing PM2.5 data from 2010 to 2014. In this data sets, its collection is based on hours, and from the original data set we can see that there are lots of

missing value (approximately 4.7% data of PM2.5 are missing). We need to delete the missing data. In order to reduce the impact of deleting missing values on data, we select the data of the fourth year and fifth year, namely 2013 and 2014 respectively.

2.2 Time Data Processing

As we have mentioned before, the PM2.5 maybe have a high co-linearity with time (a few hours a day, a few days a month, a few months a year), and in the time-series data, there usually auto-correlation between the in order to ignore this time series problem, we decided to reduce this correlation through some data processing procedures.

After setting up the full model, we discover that the R square is insignificant, then we separate the data into two classes: Summer and Autumn, Spring and Winter. We do the model estimation. Then, we classified them again and now the data are in four classes. In order to select the most appropriate categorical variable, we believe that the PM2.5 level is related to the weather. For weather data, it is usually seasonal. For example, we expect winter air pollution to become more severe as coal consumption increases. Therefore, we decided to change only the classification predictor for the seasons to categorical variables.

At the same time, the wind direction and wind force will also affect the concentration of PM2.5 every day. Since the original data set has no wind data, only the hourly wind direction, and the current data set needs the average wind direction of each day, and the average wind direction calculation. It is divided into two methods, the averaging method and the vector summation method (Yueh, 1997). In order to eliminate the "wind strength" required in the "vector summation method", we choose to use the averaging method for calculation.

In addition, through the lubridate package in the R language, we convert the character variable "time" into a numeric variable: timestamp.

In the end, the new data set contains eight independent variables and one dependent variables:

	timebyday	timestamp	season	pm2.5	Iws	Is	Ir	DEWP	TEMP	PRES	cbwd_data
1	2013-01-01	15706	Spring	16.50000	46.64667	0	0	-20.54167	-7.208333	1023.667	1.083333
2	2013-01-02	15707	Spring	18.45833	233.05625	0	0	-27.45833	-10.500000	1039.458	1.000000
3	2013-01-03	15708	Spring	24.50000	96.39458	0	0	-24.58333	-9.875000	1043.458	1.625000
4	2013-01-04	15709	Spring	78.79167	12.84167	0	0	-21.08333	-10.666667	1032.792	1.583333
5	2013-01-05	15710	Spring	66.41667	8.88625	0	0	-21.08333	-7.583333	1028.708	1.958333
6	2013-01-06	15711	Spring	131.08333	10.83833	0	0	-17.58333	-6.708333	1027.667	2.166667

FIGURE 2.3: new data set

Chapter 3

Original Data Analysis

3.1 Full Model Estimation

As we have mentioned, we did the data pre-processing before we begin our data analysis (model fitting). We analyzed this set of data according to the most fundamental steps of regression analysis. This is the analysis of the full model. Our Full model's regression formula is :
PM2.5 = $\beta_0 + \beta_1$timestamp + $\beta_2$season + β_3Iws + β_4Is + β_5Ir + $\beta_6\beta$DEWP + β_7TEMP + β_8PRESS + β_9

The model's summary are given below:

beta	Estimate	Std. Error	t value	Pr(>|t|)
(Intercept)	2.344e+03	5.414e+02	4.329	1.71e-05 ***
timestamp	-4.383e-04	1.178e-02	-0.037	0.970337
seasonSpring	2.530e+01	8.012e+00	3.158	0.001654 **
easonSummer	-3.362e+01	8.985e+00	-3.742	0.000197 ***
seasonWinter	2.408e+01	9.346e+00	2.576	0.010193 *
Iws	-1.657e-01	6.586e-02	-2.516	0.012075 *
Is	-1.544e+01	6.509e+00	-2.372	0.017953 *
Ir	-1.493e+01	2.967e+00	-5.032	6.13e-07 ***
DEWP	7.567e+00	5.021e-01	15.070	< 2e-16 ***
TEMP	-1.053e+01	7.359e-01	-14.310	< 2e-16 ***
PRES	-2.056e+00	5.292e-01	-3.885	0.000112 ***
cbwd_dataNE	-6.115e+01	1.226e+01	-4.988	7.66e-07 ***
cbwd_dataNW	-2.927e+01	7.431e+00	-3.939	9.00e-05 ***
cbwd_dataSE	-1.883e+01	6.793e+00	-2.772	0.005710 **

Which is in the form of: $y = \beta_1 x_1 + \ldots + \beta_9 x_9$

Initally, to come up with the full model, we decided to first plot correlation plots for all regressors that we believed taht may have a constant variance: See fig:Correlation of the full model

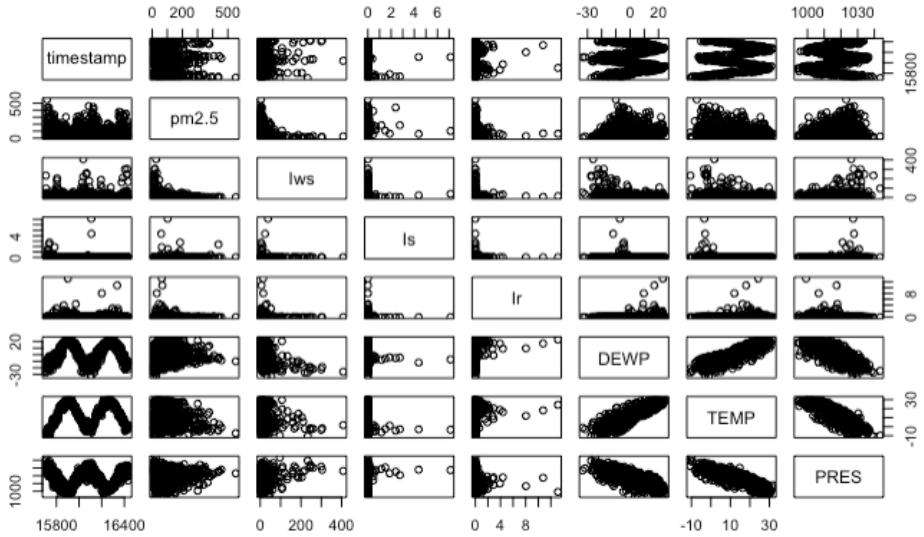

FIGURE 3.1: Correlation of the full model

From 3.1, we noticed that the correlation between PM2.5 and cumulated hour of wind speed(Iws), cumulative hours of snow(Is), and cumulative hours of rain(Ir) are all strong and their correlation density curve seem like all right skewed. We can use the 'Residuals vs Fitted' plot and the 'Normal Q-Q' plot of residual to check whether our deduction.

The '*lm*' function can also plots out the 'Residuals vs Fitted' plot, 'Normal Q-Q'plot, 'Scale-Location' and 'Residuals vs Leverage' plot: See fig:Testing Plots of Full Model

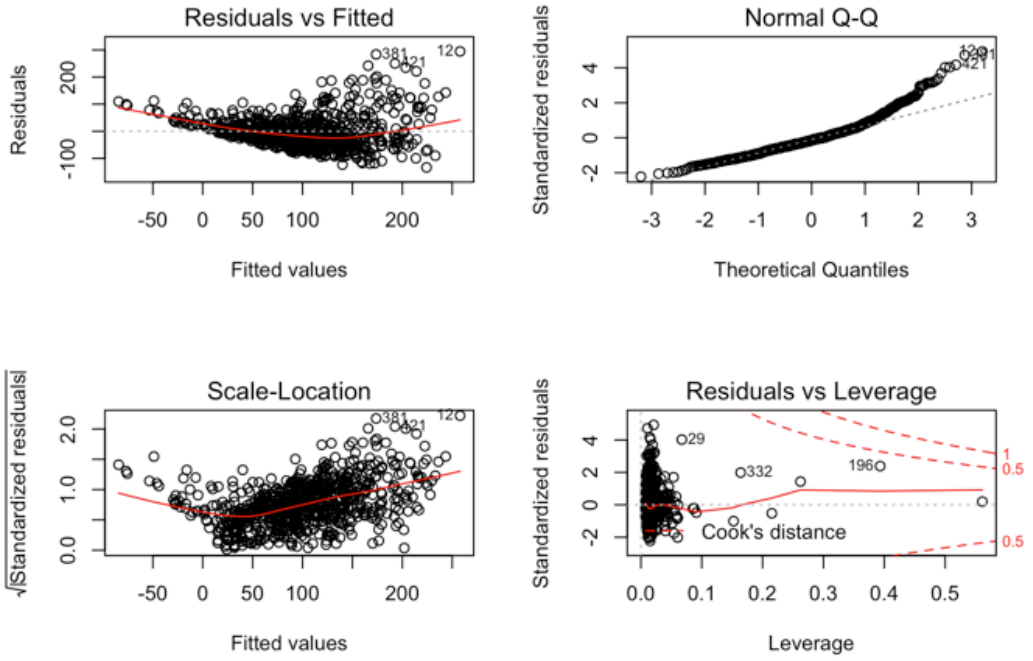

FIGURE 3.2: Four testing plots

We also use lm and vif function to check the vif value of each variables. The result

shows that all the variables vif value are not way larger than 10. The largest one is approximately 13.4, which means the variables dont have a high multicollieearity and the model has no compressibility. We cant do the ridge regression and also the lasso regression to this model.

3.2 Model with Natural Logrithm

It can be seen that the residual of the full model 3.1 is nonlinear. Usually, we need to do the **ncv test** to check want kinds of transformation should we do, however, according the reference book, we can do the Natural Logarithm transforamtion to the dependent variable when the residual plot is right skewed the(Montgomery, Peck, and Vining, 2012). Then we take the natural logarithm of the dependent variable, thus obtaining a new regression model: Full Model with Natural Logarithm Transformation, by comparing the residual graphs of the two models, we can clearly see that the residual of the model after natural logarithm transformation is linear, but we can see the residuals plot as the double-bow model(Montgomery, Peck, and Vining, 2012), which represents the model at this time. The variance of the residuals is not the same, and it does not satisfy the characteristics of the regression model of homoscedasticity.

Here is the formula of the model after the Natural Logarithm transformation: $\log(\text{PM2.5}) = \beta_0 + \beta_1 \text{timestamp} + \beta_2 \text{season} + \beta_3 \text{Iws} + \beta_4 \text{Is} + \beta_5 \text{Ir} + \beta_6 \beta \text{DEWP} + \beta_7 \text{TEMP} + \beta_8 \text{PRESS} + \beta_9$ Which is in form of : $log(y) = \beta_1 x_1 + \ldots + \beta_9 x_9$

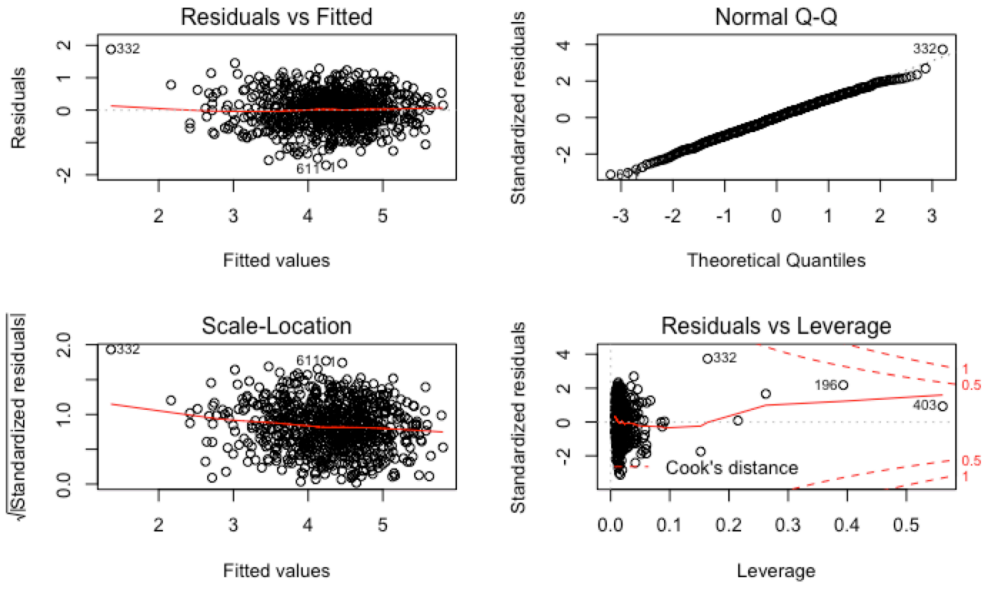

FIGURE 3.3: Full Model with Natural Logarithm test plot

3.3 Interaction Model

According to the 3.1, there are some correlations between some of the independent variables and dependent variable, It can be seen that although the residual at this time is in accordance with the normal distribution, it can be seen from the residual map that the variance of the residuals is not the same. We suspect that there are

interactions between different variables, so we have newly created a model to try to improve the R^2 of the model, and then filter the model through stepwise method.

The Interaction Model is:

$\log(\text{PM2.5}) = ()\beta_0 + \beta_1 \text{timestamp} + \beta_2 \text{season} + \beta_3 \text{Iws} + \beta_4 \text{Is} + \beta_5 \text{Ir} + \beta_6 \beta \text{DEWP} + \beta_7 \text{TEMP} + \beta_8 \text{PRESS} + \beta_9$

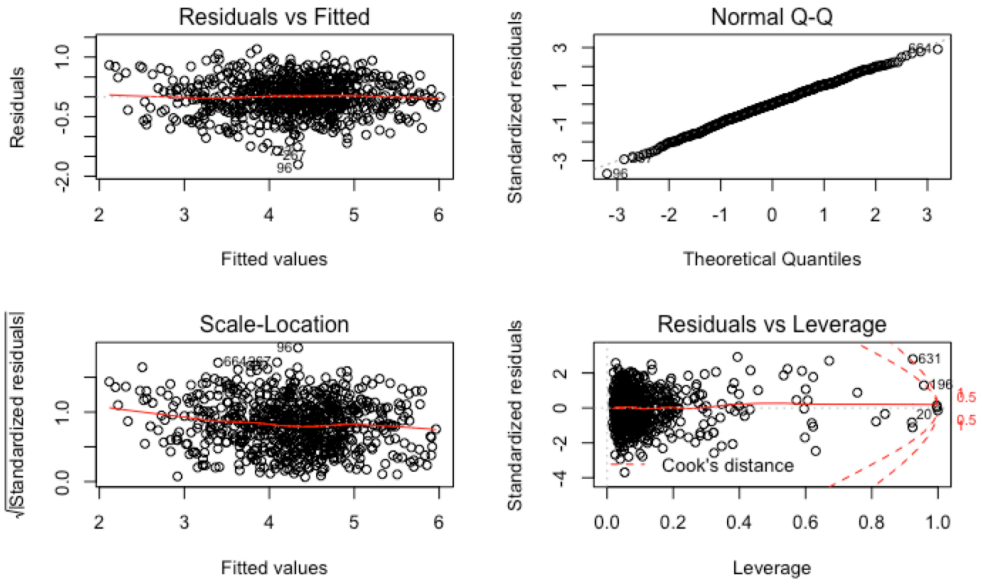

FIGURE 3.4: Natural Logarithm Model with Interaction

As we can see in the 3.4, the variance is more dispersed, the homoscedasticity is better, and R square($R^2 = 0.7086$) is higher.

3.4 Stepwise Method

In Statistics, we usually use Stepwise to choose the variables, choose the predictive varaibles by the R's built-in automatic procedure. At the beginning, the AIC = -1001.82 and at the last step, the AIC = $-1035,45$, and the model selected by AIC is: log(pm2.5) = timestamp + season + Iws + Is + Ir + DEWP + TEMP + PRES + cbwd$_d$ata + $timestamp : season + timestamp : Is + timestamp : cbwd_data + season : Iws + season : DEWP + season : TEMP + season : PRES + season : cbwd_data + Iws : Ir + Iws : DEWP + Iws : TEMP + Iws : PRES + Iws : cbwd_data + Is : TEMP + Is : PRES + Ir : DEWP + Ir : TEMP + Ir : PRES + DEWP : PRES + TEMP : cbwd_data + PRES : cbwd_data + intercept$ The R square now is 0.6995 and the test plots is much better than before, see 3.5

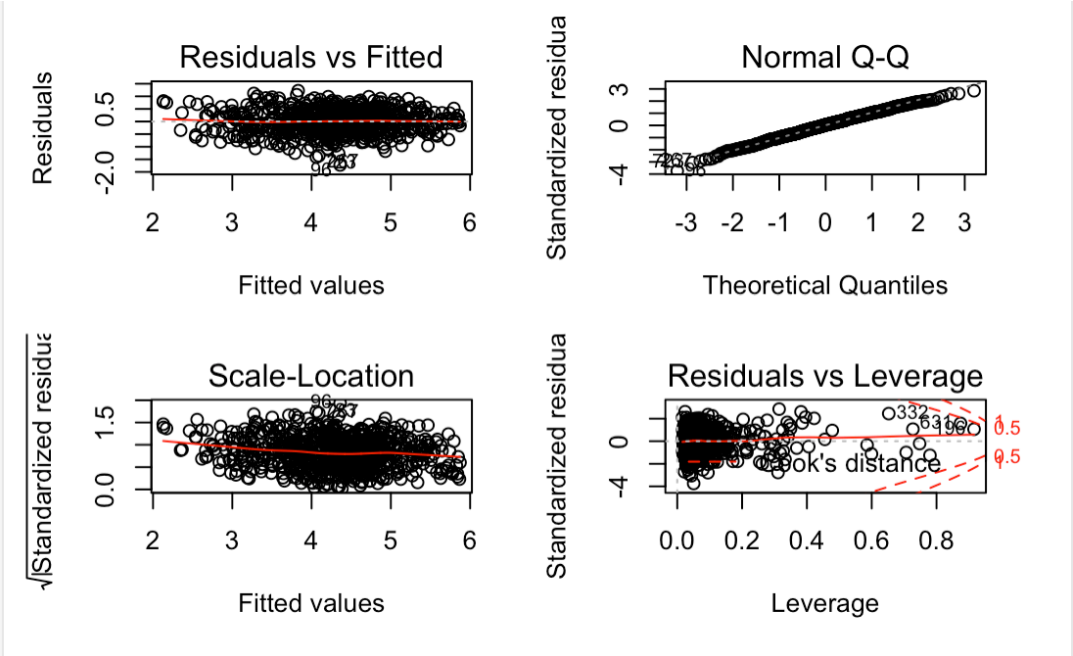

FIGURE 3.5: Stepwise Model

Although it seem like a good model, in some ways, we thought it was not a fitted model we want. After checking the original variables again, we found that the Cumulated hours of snow , Dew Point and Temperature thereT all have strong correlation with time. Then, we came up a new idea, we have to separate the time as a more smaller segments in terms of season as unit. Usually, smog is more likely to affect humans in the spring and winter seasons. This can be reflected in the previous model. The correlation between spring and winter and pm2.5 is much larger than that in summer and autumn. Then, we chose Spring + Winter as the time nodes and continue our model estimation.

Chapter 4

New Model Esitmation and Data Diagnostics

4.1 Spring Winter Model

By using the 'lm' function from the car package, we find that spring and winter are significant than two other seasons. See: fig:Coefficient Plot

```
Coefficients:
               Estimate Std. Error t value Pr(>|t|)
(Intercept)   2.344e+03  5.414e+02   4.329 1.71e-05 ***
timestamp    -4.383e-04  1.178e-02  -0.037 0.970337
seasonSpring  2.530e+01  8.012e+00   3.158 0.001654 **
seasonSummer -3.362e+01  8.985e+00  -3.742 0.000197 ***
seasonWinter  2.408e+01  9.346e+00   2.576 0.010193 *
Iws          -1.657e-01  6.586e-02  -2.516 0.012075 *
Is           -1.544e+01  6.509e+00  -2.372 0.017953 *
```

FIGURE 4.1: Coefficient Plot

So, we decide to grouping the data into two parts. Spring-Winter data as one set and the Summer-Autumn data as another set. Then we use the lm function to compare these two models and find that the Spring-Winter Model is better than the other one.

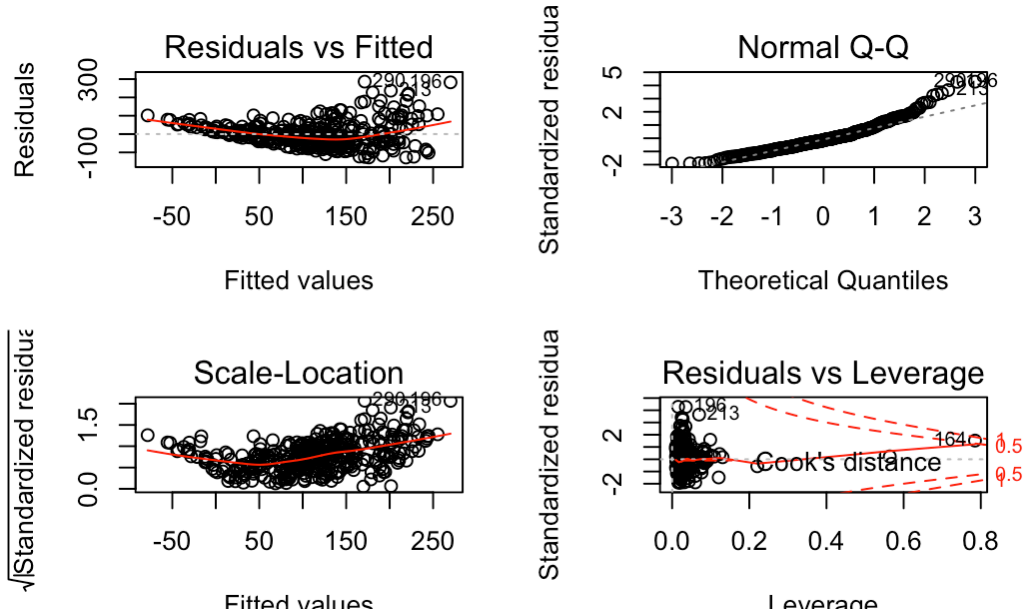

FIGURE 4.2: Spring Winter Model

From the plots we can see that there is not linear correlation between these variables. We thus do the ncvTest and draw the spreadlevel plot, the suggested power transformation is 0.381492.Kabacoff2015 See: fig:NCVTest

```
Non-constant Variance Score Test
Variance formula: ~ fitted.values
Chisquare = 104.9767, Df = 1, p = < 2.22e-16
```

FIGURE 4.3: NCV Test

```
Non-constant Variance Score Test
Variance formula: ~ fitted.values
Chisquare = 104.9767, Df = 1, p = < 2.22e-16

22 negative fitted values removed
Suggested power transformation:  0.381492
```

FIGURE 4.4: NCV Test and Suggested Power Transformation

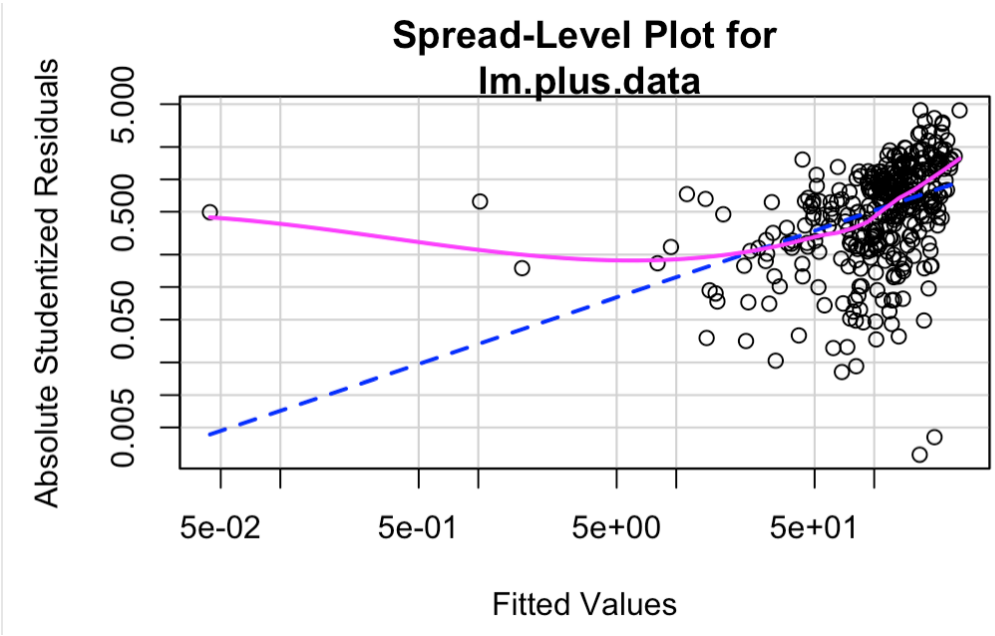

FIGURE 4.5: Spreadlevel Plot

. We think it may means closed to 0.5 or closed to 0. First, we take square loot to the pm2.5, but the result ignores our consideration. Then we start to consider another possibility, which is closed to 0, meaning we should take the natural logarithm to the pm2.5. We wonder what will happen if we take the natural logarithm to the pm2.5. We use lm function again to estimate the Spring-Winter Model with natural logarithm and the results shows the linear correlation. Also, the residual plots and the R square are both better than before.

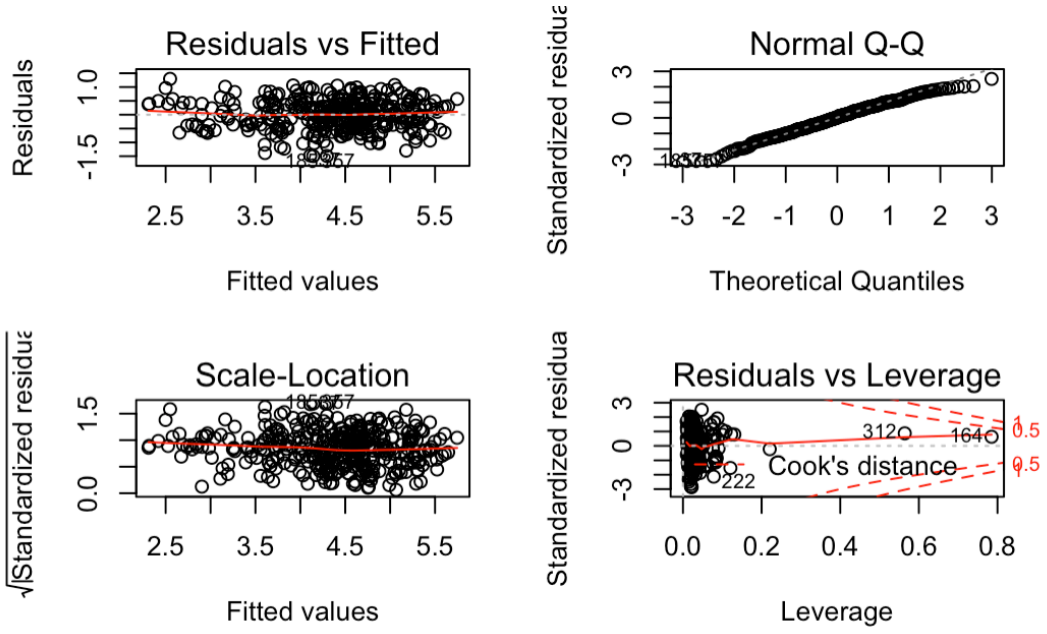

FIGURE 4.6: Natural Logarithm of Spring Winter Model

From the full model vif figure, the vif value of the winter is the most significant one in four single seasons. With this discovery, we choose the winter data to form

a model with natural logarithm. Also, we form a model for the spring data with natural logarithm, but it is not good as the winter one.

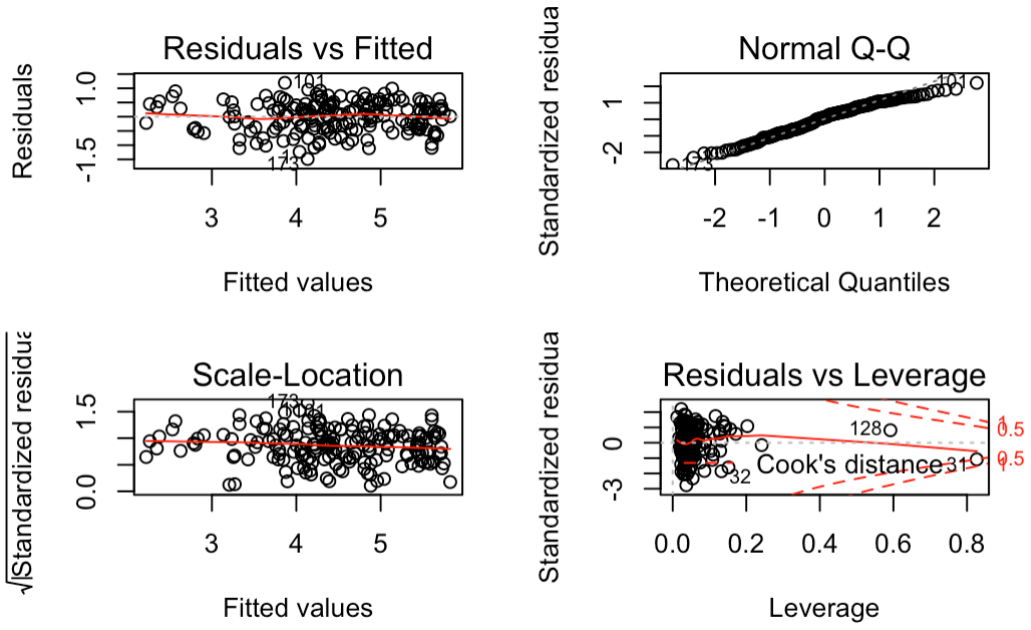

FIGURE 4.7: Winter Model Natural Logarithm

The Formula of Winter Model with Natural Logarithm is :

$\log(\text{pm2.5}) = \beta_1 timestamp + \beta_2 Iws + \beta_3 Is + \beta_4 Ir + \beta_5 DEWP + \beta_6 TEMP + \beta_7 PRES + \beta_8 cbwd_data$

The residual plots are better and the R square is higher than what we have found out before. So we use this model 4.1 to do the regression model diagnostics.

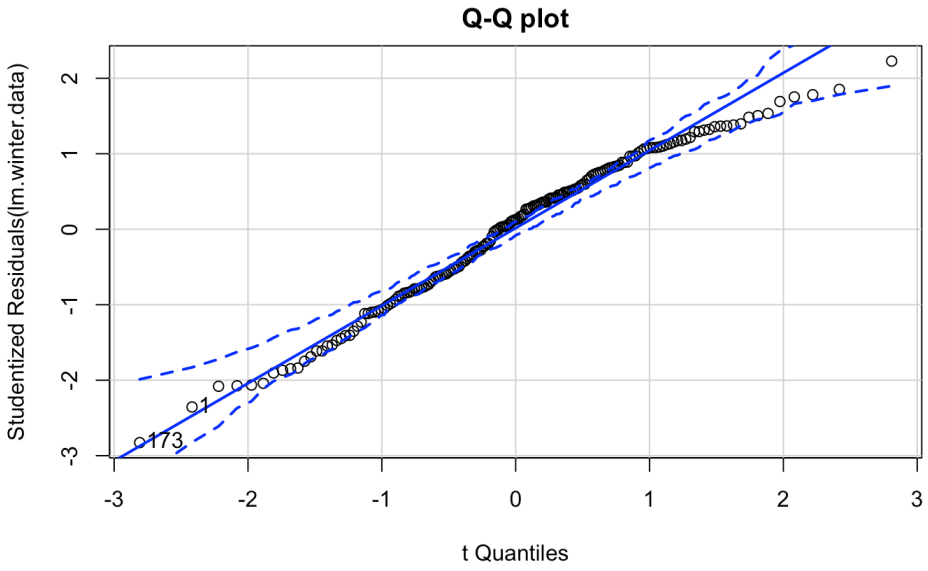

FIGURE 4.8: Q-Q Plot

First, we draw the Q-Q plot, the plot shows that all the points are closed to the straight line, and they are all in the confident interval, which means the normality assumption of this Winter Natural Logarithm Model is good. We also use the residplot function to draw the Studentized Residual Histogram, and add the Normal Curve, Kernel Density Curve and Rug Plot.

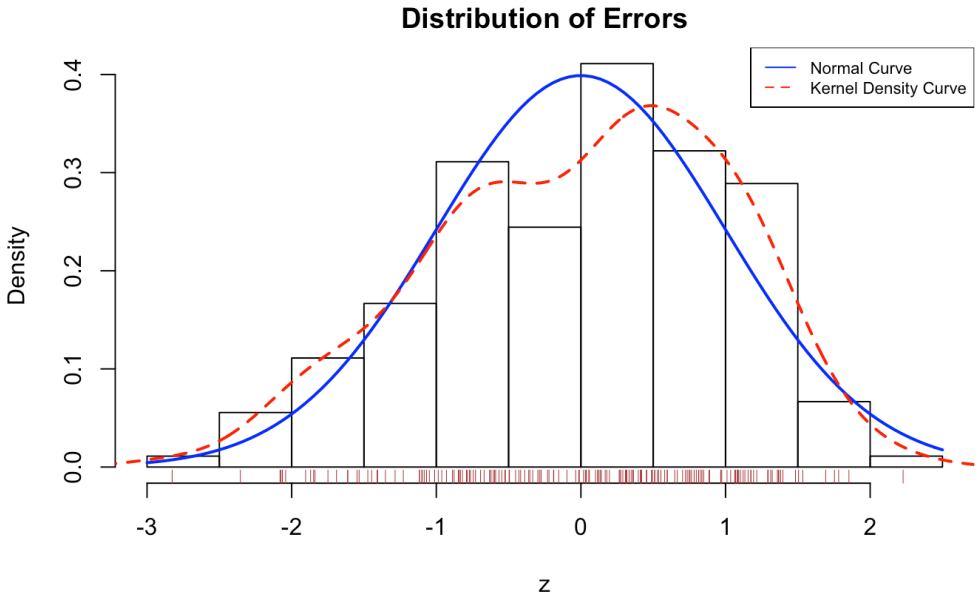

FIGURE 4.9: Studentized Residual Density Plot

Chapter 5

Conclusion

5.1 Conlcusion

Actually, we do not have a conclusion, we just have some discovery. After we try lots of estimate methods, such as taking natural logarithm, using interaction, stepwise method, ridge and lasso regression, we still cannot estimate the data in a right way. So we think if we want to use linear regression methods to estimate our data, which is kind of time series data. Maybe we should divide the data into several single segments, then do the linear regression estimation to each single segments.

Appendix A

R Code

The color of links can be changed to your liking using:

```
  library(VIM) # function aggr: visualize the missing value
library(tidyverse) #To use ggplot2, tidyr, dplyr
library(plotly) #To create interactive plots
library(DT) #To display the data
library(magrittr) #To pipe operators
library(ggplot2) #To make and customize quickly plots
library(devtools) #To Make Developing R Packages Easier
library(lubridate) # date tranformation
library(beginr)

beijing.data <- read.csv("PRSA_data_2010.1.1-2014.12.31.csv", header = T) # load the data
head(beijing.data)
tail(beijing.data)

sum(is.na(beijing.data))
aggr(beijing.data, prop = T, number = T)

i <- NULL
j <- 1
compare_value_j <- 1
for ( i in 2010:2014){
  data_i <- beijing.data[beijing.data$year == i,]
  if (compare_value_j < length(na.omit(data_i$pm2.5))){
    compare_value_j <- length(na.omit(data_i$pm2.5))
    j <- j + 1
  }
  print(j + 2009) # the year will least missing value
}
beijing.data <- as_tibble(beijing.data)

i <- NULL
for ( i in 1:length(beijing.data$No)){
  if(beijing.data$month[i] == 3){
    beijing.data$season[i] <- 1
  }
  if(beijing.data$month[i] == 4){
    beijing.data$season[i] <- 1
  }
  if(beijing.data$month[i] == 5){
```

```
    beijing.data$season[i] <- 1
  }
  if(beijing.data$month[i] == 6){
    beijing.data$season[i] <- 2
  }
  if(beijing.data$month[i] == 7){
    beijing.data$season[i] <- 2
  }
  if(beijing.data$month[i] == 8){
    beijing.data$season[i] <- 2
  }
  if(beijing.data$month[i] == 9){
    beijing.data$season[i] <- 3
  }
  if(beijing.data$month[i] == 10){
    beijing.data$season[i] <- 3
  }
  if(beijing.data$month[i] == 11){
    beijing.data$season[i] <- 3
  }
  if(beijing.data$month[i] == 12){
    beijing.data$season[i] <- 4
  }
  if(beijing.data$month[i] == 1){
    beijing.data$season[i] <- 4
  }
  if(beijing.data$month[i] == 2){
    beijing.data$season[i] <- 4
  }
}
head(beijing.data)

cleanbeijing <-select(beijing.data, c("year","month","day","hour","season","pm2.5","c
  na.omit() %>%
  filter(year >= 2013)%>%
  unite(timebyday, c("year", "month", "day"), remove = FALSE, sep = "-")
datatable(cleanbeijing, option = list(scrollX = TRUE))

#calculate the PM2.5 by day
daypm<-cleanbeijing%>%
  group_by(timebyday)%>%
  summarise(mean=mean(cleanbeijing$pm2.5))%>%
  as_tibble()
#calculate the PM2.5 by year
cleanbeijing$quality <- ifelse(cleanbeijing$pm2.5 <= 50, "good",
                               ifelse(cleanbeijing$pm2.5 <= 100, "moderate",
                                      ifelse(cleanbeijing$pm2.5 <= 300, "unhealthy"
qualitypm <- cleanbeijing %>%
  group_by(year, quality) %>%
  count() %>%
  as_tibble()
```

```
ggplot(qualitypm, aes(x = factor(year) , y = n,fill = quality)) + geom_bar(stat = 'identi
  theme(legend.title = element_blank())


spring<-filter(cleanbeijing,cleanbeijing$season==1)
summer<-filter(cleanbeijing,cleanbeijing$season==2)
autumn<-filter(cleanbeijing,cleanbeijing$season==3)
winter<-filter(cleanbeijing,cleanbeijing$season==4)

seasonpm<- cleanbeijing %>%
  group_by(season,quality)%>%
  count()%>%
  as_tibble()

ggplot(seasonpm, aes(x = factor(season) , y = n,fill = quality)) + geom_bar(stat = 'ident
  theme(legend.title = element_blank())

  cleanbeijing <- as.data.frame(cleanbeijing)
cleanbeijing <- cleanbeijing[,-c(2,3,4)]
head(cleanbeijing)
time <- cleanbeijing$timebyday
time <- as.Date(as.POSIXct(ymd(time), origin = "2013-01-01"))
cleanbeijing$timebyday <- time
cleanbeijing$timestamp <- as.numeric(cleanbeijing$timebyday)
head(cleanbeijing)
tail(cleanbeijing)
#
# cleanbeijing <- cleanbeijing[,-c(2,3,4,5)]

for (i in 1:length(cleanbeijing$timebyday)){
  if(cleanbeijing$cbwd[i] == "NW"){
    cleanbeijing$cbwd_data[i] = 1
  }
  if(cleanbeijing$cbwd[i] == "cv"){
    cleanbeijing$cbwd_data[i] = 2
  }
  if(cleanbeijing$cbwd[i] == "NE"){
    cleanbeijing$cbwd_data[i] = 3
  }
  if(cleanbeijing$cbwd[i] == "SE"){
    cleanbeijing$cbwd_data[i] = 4
  }
}

cleanbeijing_combin <- tapplydf(cleanbeijing, c("timestamp","season", "pm2.5", "Iws", "Is

FindMode <- function(x) {
    ux <- unique(x)
    ux[which.max(tabulate(match(x, ux)))]
}
```

```
# cleanbeijing_combin$cbwd_data <- rep(1,length(cleanbeijing_combin$timebyday))


cleanbeijing_combin$cbwd_data <-  tapply(cleanbeijing$cbwd_data, cleanbeijing$timebyd
cleanbeijing_combin$cbwd_data <- as.numeric(cleanbeijing_combin$cbwd_data)
# cleanbeijing_combin <- cleanbeijing_combin[,-1]
# lm.cleanbeijing_combin <- lm(pm2.5~., data = cleanbeijing_combin)
# summary(lm.cleanbeijing_combin)

i <- 1
for ( i in 1:length(cleanbeijing_combin$timestamp)){
    if(month(cleanbeijing_combin$timebyday)[i] == 3){
        cleanbeijing_combin$season[i] <- "Spring"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 4){
        cleanbeijing_combin$season[i] <- "Spring"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 5){
        cleanbeijing_combin$season[i] <- "Spring"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 6){
        cleanbeijing_combin$season[i] <- "Summer"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 7){
        cleanbeijing_combin$season[i] <- "Summer"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 8){
        cleanbeijing_combin$season[i] <- "Summer"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 9){
        cleanbeijing_combin$season[i] <- "Autumn"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 10){
        cleanbeijing_combin$season[i] <- "Autumn"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 11){
        cleanbeijing_combin$season[i] <- "Autumn"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 12){
        cleanbeijing_combin$season[i] <- "Winter"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 1){
        cleanbeijing_combin$season[i] <- "Winter"
    }
    if(month(cleanbeijing_combin$timebyday)[i] == 2){
        cleanbeijing_combin$season[i] <- "Winter"
    }
}
head(cleanbeijing_combin)

# cleanbeijing_combin$cbwd_data <- floor(cleanbeijing_combin$cbwd_data)
```

```
for (i in 1:length(cleanbeijing_combin$timebyday)){
  if(cleanbeijing_combin$cbwd[i] == 1){
    cleanbeijing_combin$cbwd_data[i] = "NW"
  }
  if(cleanbeijing_combin$cbwd_data[i] == 2){
    cleanbeijing_combin$cbwd_data[i] = "cv"
  }
  if(cleanbeijing_combin$cbwd_data[i] == 3){
    cleanbeijing_combin$cbwd_data[i] = "NE"
  }
  if(cleanbeijing_combin$cbwd_data[i] == 4){
    cleanbeijing_combin$cbwd_data[i] = "SE"
  }
}

lm.cleanbeijing_combin.nolog <- lm(pm2.5~timestamp + season +Iws + Is + Ir + DEWP + TEMP
summary(lm.cleanbeijing_combin.nolog)
par(mfrow=c(2,2))
plot(lm.cleanbeijing_combin.nolog)

library(car)
a <- as.data.frame(cleanbeijing_combin[,-c(1,3,11)])
cor(a)
pairs(a)

lm.cleanbeijing_combin <- lm(log(pm2.5)~timestamp + season +Iws + Is + Ir + DEWP + TEMP +
summary(lm.cleanbeijing_combin)
par(mfrow=c(2,2))
plot(lm.cleanbeijing_combin)

lm.cleanbeijing_combin_interaction <- lm(log(pm2.5)~(timestamp + season + Iws + Is + Ir +
summary(lm.cleanbeijing_combin_interaction)
par(mfrow=c(2,2))
plot(lm.cleanbeijing_combin_interaction)

lm.cleanbeijing_combin_step <- step(lm.cleanbeijing_combin_interaction, direction = "both
summary(lm.cleanbeijing_combin_step)
par(mfrow=c(2,2))
plot(lm.cleanbeijing_combin_step)

winter_data<-filter(cleanbeijing_combin,cleanbeijing_combin$season=="Winter")
names(winter_data)
lm.winter.data <- lm(log(pm2.5)~timestamp + Iws + Is + Ir + DEWP + TEMP + PRES + cbwd_dat
summary(lm.winter.data)
par(mfrow=c(2,2))
plot(lm.winter.data)

Spring_data<-filter(cleanbeijing_combin,cleanbeijing_combin$season=="Spring")
names(Spring_data)
lm.spring.data <- lm(log(pm2.5)~timestamp + Iws + Is + Ir + DEWP + TEMP + PRES + cbwd_dat
summary(lm.spring.data)
```

```
par(mfrow=c(2,2))
plot(lm.spring.data)

spring_winter <- as.data.frame(rbind(Spring_data, winter_data))
names(spring_winter)

for (i in 1:length(spring_winter$timebyday)){
  if(spring_winter$season == 1){
    spring_winter$season[i] = "Spring"
  }
  if(spring_winter$season == 4){
      spring_winter$season[i] = "Winter"
    }
}
lm.plus.data <- lm(log(pm2.5)~timestamp + Iws + Is + Ir + DEWP + TEMP + PRES + cbwd_d
summary(lm.plus.data)
par(mfrow=c(2,2))
plot(lm.plus.data)

library(car)
qqPlot(lm.winter.data, labels = row.names(winter_data), id.methods = "identify", simu

residplot <- function(fit, nbreaks = 10){
  z <- rstudent(fit)
  hist(z, breaks = nbreaks, freq = FALSE,
       xlib = "Studentized Residual",
       main = "Distribution of Errors")
  rug(jitter(z), col = "brown")
  curve(dnorm(x), mean = mean(z), sd = sd(z), add = TRUE, col = "blue",lwd = 2)
lines(density(z)$x, density(z)$y,
      col="red", lwd = 2, lty = 2)
legend("topright",
       legend = c("Normal Curve", "Kernel Density Curve"),
       lty = 1:2, col = c("blue", "red"), cex=.7)
}
residplot(lm.winter.data)

library(car)
ncvTest(lm.winter.data)
spreadLevelPlot(lm.plus.data)
```

Bibliography

Liang, Xuan et al. (2015). "Assessing Beijing's PM2. 5 pollution: severity, weather impact, APEC and winter heating". In: *Proc. R. Soc. A* 471.2182, p. 20150257.

Montgomery, Douglas C, Elizabeth A Peck, and G Geoffrey Vining (2012). *Introduction to linear regression analysis*. Vol. 821. John Wiley & Sons.

Yueh, Simon H (1997). "Modeling of wind direction signals in polarimetric sea surface brightness temperatures". In: *IEEE Transactions on Geoscience and Remote Sensing* 35.6, pp. 1400–1418.